

Want to run your data access layer at warp speed?

Your email address..

Subscribe



## VLAD MIHALCEA

High-Performance Java Persistence and Hibernate



# The anatomy of Hibernate dirty checking mechanism

AUGUST 21, 2014 / VLADMIHALCEA


*(Last Updated On: January 29, 2018)*

## Introduction

The persistence context enqueues **entity state transitions** that get translated to database statements upon flushing. For managed entities, Hibernate can auto-detect incoming changes and schedule SQL UPDATES on our behalf. This mechanism is called **automatic dirty checking**.

## The default dirty checking strategy

By default Hibernate checks all managed entity properties. Every time an entity is loaded, Hibernate makes an additional copy of all entity property values. At flush time, every managed entity property is matched against the loading-time snapshot value:



So the number of individual **dirty checks** is given by the following formula:

$$N = \sum_{k=1}^n p_k$$

where

n = The number of managed entities

p = The number of properties of a given entity

Even if only one property of a single entity has ever changed, Hibernate will still check all managed entities. For a large number of managed entities, the default dirty checking mechanism may have a significant CPU and memory footprint. Since the initial entity snapshot is held separately, the persistence context requires twice as much memory as all managed entities would normally occupy.

# Bytecode instrumentation

A more efficient approach would be to mark dirty properties upon value changing. Analogue to the original deep comparison strategy, it's good practice to decouple the domain model structures from the change detection logic. The automatic entity change detection mechanism is a **cross-cutting concern**, that can be woven either at build-time or at runtime.

The entity class can be appended with bytecode level instructions implementing the automatic dirty checking mechanism.

## Weaving types

The bytecode enhancement can happen at:

- Build-time

After the hibernate entities are compiled, the build tool (e.g. **ANT**, **Maven**) will insert bytecode level instructions into each compiled entity class. Because the classes are enhanced at build-time, this process exhibits no extra runtime penalty. Testing can be done against enhanced class versions, so that the actual production code is validated before the project gets built.

- Runtime

The runtime weaving can be done using:

- A Java agent, doing bytecode enhancement upon entity class loading
- A runtime container (e.g. **Spring**), using **JDK Instrumentation support**

If you enjoyed this article, I bet you are going to love my **Book** and **Video Courses** as well.



## Hibernate 5 improvements

Hibernate 3 has been offering bytecode instrumentation through **an ANT target** but it never became mainstream and most Hibernate projects are still currently using the default deep comparison approach.

**Hibernate 5** has redesigned the **bytecode enhancement mechanism**, is more reliable than it used to be.

## Subscribe to our Newsletter

\* indicates required

Email Address \*

**10 000** readers have found this blog worth following!

If you **subscribe** to my newsletter, you'll get:

- A **free sample** of my Video Course about running Integration tests at warp-speed using Docker and tmpfs

- **3 chapters** from my book, **High-Performance Java Persistence**,
- a **10% discount** coupon for my book.

Get the most out of your persistence layer!

Subscribe

Advertisements

---

**Related**

How to enable bytecode enhancement dirty checking in Hibernate  
February 11, 2016  
In "Hibernate"

How to customize Hibernate dirty checking mechanism  
August 29, 2014  
In "Hibernate"

How to prevent lost updates in long conversations  
September 22, 2014  
In "Hibernate"

Categories: [Hibernate](#), [Java](#)    Tags: [automatic dirty checking](#), [hibernate](#), [Session flush](#), [Training](#), [Tutorial](#)

[← How does AUTO flush strategy work in JPA and Hibernate](#)

[How to customize Hibernate dirty checking mechanism →](#)

---

## Leave a Reply

Your email address will not be published. Required fields are marked \*

**Comment**

Name \*

Email \*

Website


☐ Save my name, email, and website in this browser for the next time I comment.

☐ Sign me up for the newsletter!

Post Comment


☐ Notify me of follow-up comments by email.

## VIDEO COURSE




## HIGH-PERFORMANCE JAVA PERSISTENCE





WIN A COPY OF MY BOOK!



## HIBERNATE PERFORMANCE TUNING TIPS


HYPERSTISTENCE




WIN A COPY OF MY BOOK!



HIGH-PERFORMANCE JAVA PERSISTENCE



## VIDEO COURSE



Search ...



[RSS - Posts](#)

[RSS - Comments](#)

## ABOUT

[About](#)

[Privacy Policy](#)

[Terms of Service](#)

**Download free chapters and get a 10% discount for the "High-Performance Java Persistence" book**

