

The Most Efficient Data Type For To-Many Associations

The difference between a Bag and a List

Hibernate's naming of the different collection types is a little bit confusing because *Lists* and *Bags* are both mapped by a *java.util.List*. The difference between them is that a *List* is ordered and a *Bag* is unordered.

You should prefer the unordered *Bag* for most mappings because retrieving the association in a specific order slows down your database queries. You should better use a [JPQL query](#) with an ORDER BY clause to define the ordering if you need it.

Should you use a Bag or a Set?

When you just look at the Java types, the answer seems to be easy. In general, a *java.util.List* provides the better performance while a *java.util.Set* doesn't contain any duplicates. As long as you implement the create use case correctly, a *java.util.List* seems like the obvious best choice for your association mapping.

But it's not that easy. A *List* might be more efficient than a *Set*, but the type also influences how Hibernate manages the association in the database. So, there are a few other things you need to take into account when you make your decision.

A critical bug in older Hibernate versions

If you're using a Hibernate version older than 5.0.8, you should be aware of bug [HHH-5855](#). When you used a *java.util.List* and [merged the parent entity](#), Hibernate generated 2 INSERT statements for each new child entity.

The Most Efficient Data Type For To-Many Associations

Inefficient handling of many-to-many associations

When you're mapping a many-to-many association, you should always use a `java.util.Set`.

```
@Entity
public class Book {

    @ManyToMany
    @JoinTable(name = "book_author",
        joinColumns = { @JoinColumn(name = "fk_book") },
        inverseJoinColumns = { @JoinColumn(name = "fk_author") })
    private Set authors = new HashSet();

    ...
}
```

If you model the association as a `java.util.List`, Hibernate handles the removal of associated entities very inefficiently. Instead of removing the record that maps the removed association, Hibernate removes all records from the association table before it inserts a new record for the remaining associations.

Summary

As you've seen, mapping an association as a *java.util.List* can create problems which by far outweigh the small performance gain you get compared to a *java.util.Set*. So, better make sure to update your Hibernate version and to use a *Set* to model many-to-many associations.

